Mind the Cost of Telemetry Data Analysis

Alessandra Fais[1], Gianni Antichi[2], Stefano Giordano[1], Giuseppe Lettieri[1], Gregorio Procissi[1]

[1] Università di Pisa, Italy | [2] Queen Mary University of London, United Kingdom

Introduction

- Stream processing engines **efficiently process continuous amounts (streams) of information**
  - Widely used solutions ( ) for a variety of use cases
- Network operators need efficient ways to analyze fine-grained telemetry data
  - In production datacenter networks, hundreds of thousands of switches produce up to millions of reports per second!

GOAL: What's the best streaming engine for network traffic analysis?

A Qualitative Comparison

	Flink	Spark	Storm	WindFlow
Batching	✗	mandatory	✗	✓
Chaining	✓	✗	✗	✓
Ordering	✗	btw batches	✗	✓
Windows	✓	✓	✗	✓
Event time	✓	partial	✗	✓
Distributed	✓	✓	✓	✗

Heavy hitters → dst hosts hit by at least one heavy flow in a time window

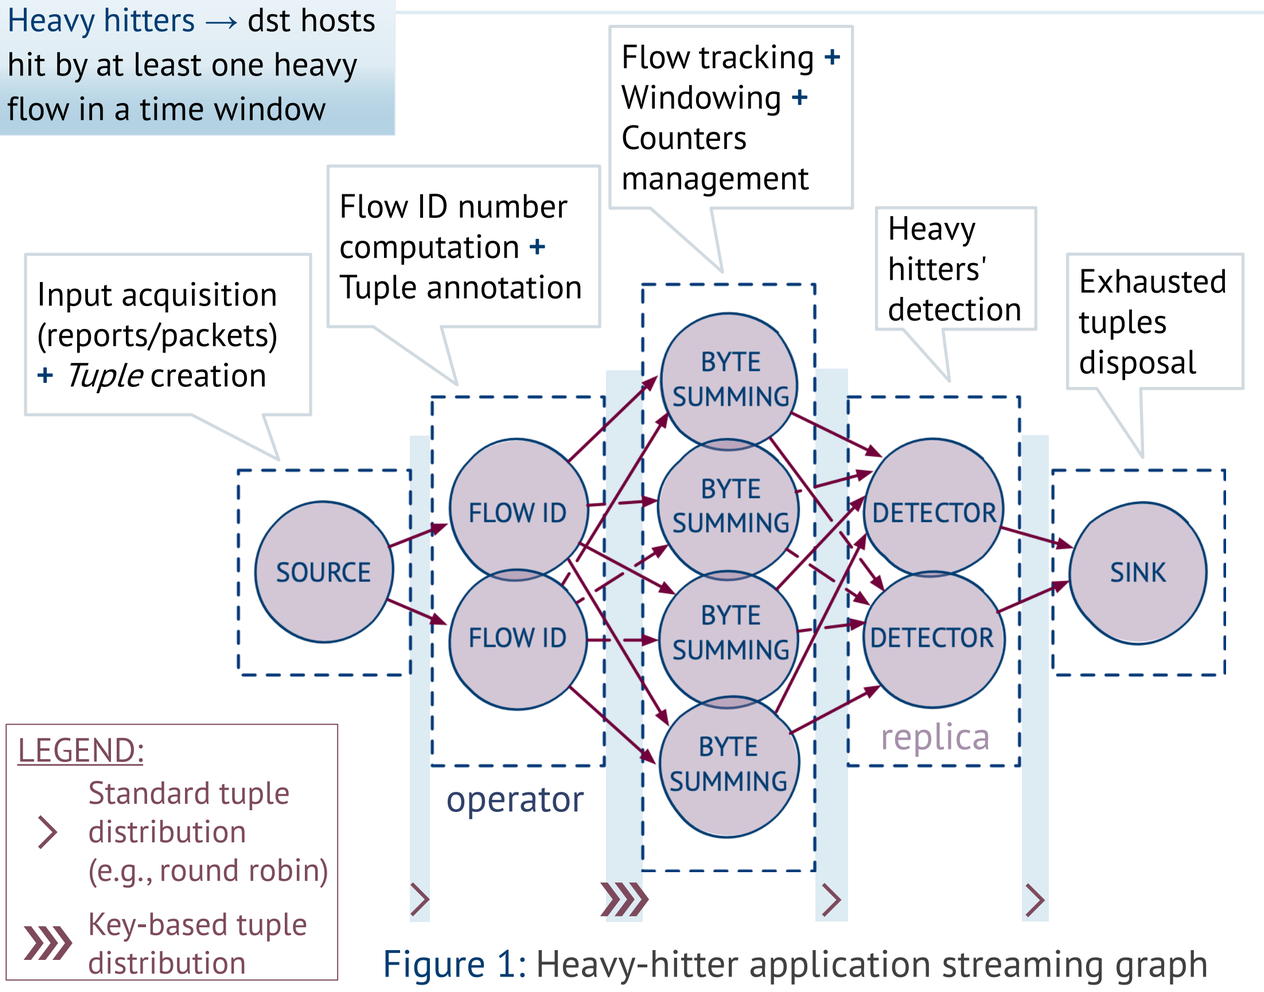

Figure 1: Heavy-hitter application streaming graph

Findings

- Systems designed for generic data processing over distributed platforms perform poorly with network data
  - Overheads not compensated by the computational burden of the application itself

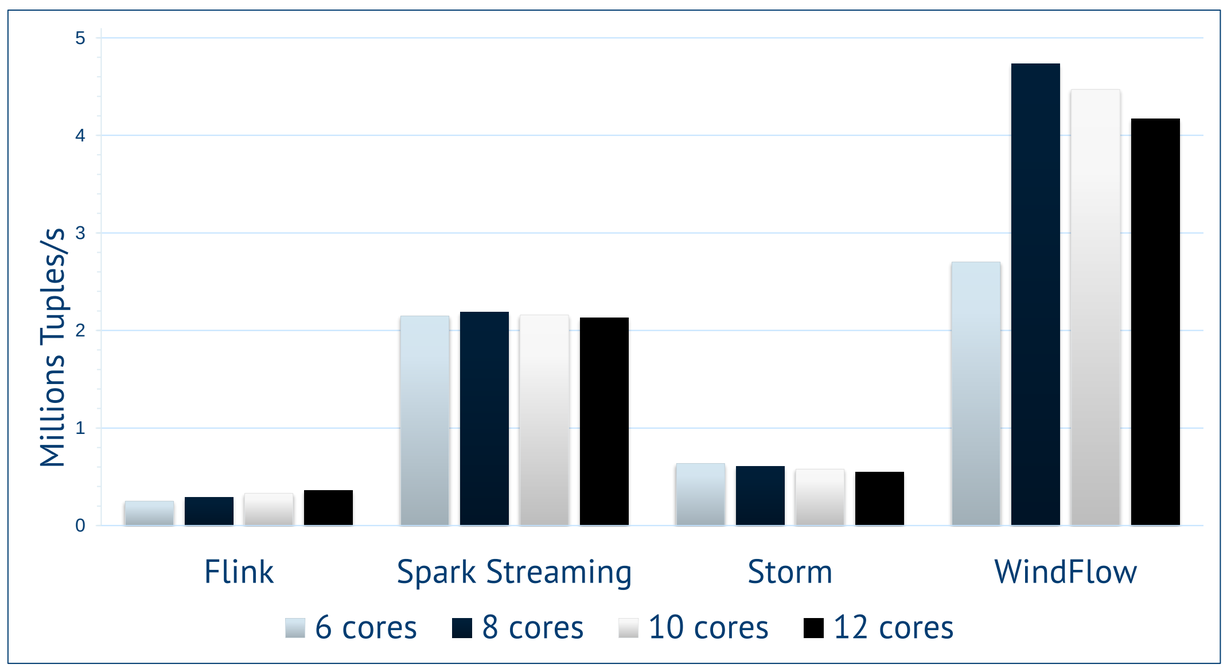

- WindFlow shows better performance figures
  - More than 2x in most cases than Spark Streaming
  - Around 10x of Flink and Storm
- WindFlow performance scales with n. physical cores
- Other solutions immediately saturate

Resource utilization scenarios
- Physical cores only → number of cores ⩽ 8
- Hyperthreading → 8 < number of cores ⩽ 16

Figure 2: Performance comparison

Promising research directions

Design of a networking domain specific streaming engine

Lightweight communication mechanisms → Support for computation distribution over a cluster → Specifically built for network traffic analysis

w/o compromising performance!

Network data analysis traits: sustained input rate + moderate comp. burden

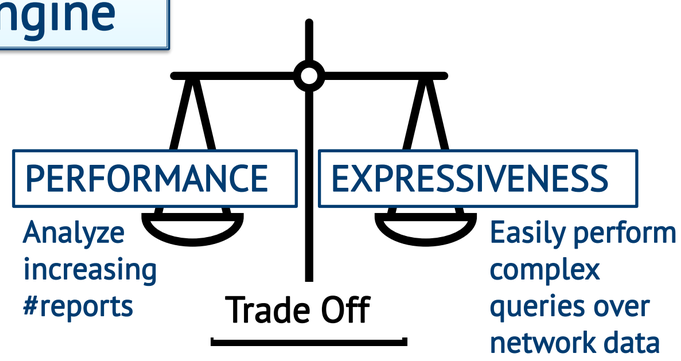

Contact information: alessandra.fais@phd.unipi.it | https://for.unipi.it/alessandra_fais/